

### **1. customer\_info.csv**

1,John,101

2,Alice,102

3,Bob,103

### **2. order\_info.csv**

101,501,2

102,502,1

103,503,5

### **3. item\_info.csv**

501,Laptop,60000

502,Mouse,500

503,Keyboard,1500

**Save this in /home/cloudera**

**[cloudera@quickstart Desktop]\$ hive**

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

WARNING: Hive CLI is deprecated and migration to Beeline is recommended.

**hive>**

**> -- Customer Information Table**

**> CREATE TABLE Customer\_info (**

**> Cust\_ID INT,**

**> Cust\_Name STRING,**

**> OrderID INT**

**> )**

**> ROW FORMAT DELIMITED**

**> FIELDS TERMINATED BY ','**

**> STORED AS TEXTFILE;**

OK

Time taken: 0.728 seconds

**hive>**

**> -- Order Info Table**

```
> CREATE TABLE Order_info (  
  >   OrderID INT,  
  >   ItemID INT,  
  >   Quantity INT  
  > )  
  
> ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE;
```

OK

Time taken: 0.093 seconds

hive>

```
> -- Item Info Table  
  
> CREATE TABLE Item_info (  
  >   Item_ID INT,  
  >   Item_Name STRING,  
  >   ItemPrice FLOAT  
  > )  
  
> ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE;
```

OK

Time taken: 0.252 seconds

hive> LOAD DATA LOCAL INPATH '/home/cloudera/customer\_info.csv' INTO TABLE Customer\_info;

Loading data to table default.customer\_info

Table default.customer\_info stats: [numFiles=1, totalSize=34]

OK

Time taken: 0.736 seconds

hive> LOAD DATA LOCAL INPATH '/home/cloudera/order\_info.csv' INTO TABLE Order\_info;

Loading data to table default.order\_info

Table default.order\_info stats: [numFiles=1, totalSize=31]

OK

Time taken: 0.417 seconds

**hive> LOAD DATA LOCAL INPATH '/home/cloudera/item\_info.csv' INTO TABLE Item\_info;**

Loading data to table default.item\_info

Table default.item\_info stats: [numFiles=1, totalSize=50]

OK

Time taken: 0.638 seconds

**hive> SELECT \* FROM Customer\_info;**

OK

1	John	101
---	------	-----

2	Alice	102
---	-------	-----

3	Bob	103
---	-----	-----

NULL	NULL	NULL
------	------	------

Time taken: 0.441 seconds, Fetched: 4 row(s)

**hive> SELECT \* FROM Order\_info;**

OK

101	501	2
-----	-----	---

102	502	1
-----	-----	---

103	503	5
-----	-----	---

NULL	NULL	NULL
------	------	------

Time taken: 0.064 seconds, Fetched: 4 row(s)

**hive> SELECT \* FROM Item\_info;**

OK

501	Laptop	60000.0
-----	--------	---------

502	Mouse	500.0
-----	-------	-------

503	Keyboard	1500.0
-----	----------	--------

NULL	NULL	NULL
------	------	------

Time taken: 0.074 seconds, Fetched: 4 row(s)

**hive> SELECT**

**> c.Cust\_ID,**

**> c.Cust\_Name,**

**> o.OrderID,**

```

> i.Item_Name,
> o.Quantity,
> i.ItemPrice,
> o.Quantity * i.ItemPrice AS TotalCost
> FROM
> Customer_info c
> JOIN
> Order_info o ON c.OrderID = o.OrderID
> JOIN
> Item_info i ON o.ItemID = i.Item_ID;

```

Query ID = cloudera\_20250507073939\_4639ce79-2d14-4772-ba51-6733d20f3593

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera\_20250507073939\_4639ce79-2d14-4772-ba51-6733d20f3593.log

2025-05-07 07:39:34 Starting to launch local task to process map join;maximum memory = 932184064

2025-05-07 07:39:36 Dump the side-table for tag: 1 with group count: 3 into file:  
file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-39-28\_145\_8930674938168841766-1/-local-10005/HashTable-Stage-5/MapJoin-mapfile01--.hashtable

2025-05-07 07:39:36 Uploaded 1 File to: file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-39-28\_145\_8930674938168841766-1/-local-10005/HashTable-Stage-5/MapJoin-mapfile01--.hashtable (357 bytes)

2025-05-07 07:39:36 Dump the side-table for tag: 1 with group count: 3 into file:  
file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-39-28\_145\_8930674938168841766-1/-local-10005/HashTable-Stage-5/MapJoin-mapfile11--.hashtable

2025-05-07 07:39:36 Uploaded 1 File to: file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-39-28\_145\_8930674938168841766-1/-local-10005/HashTable-Stage-5/MapJoin-mapfile11--.hashtable (329 bytes)

2025-05-07 07:39:36 End of local task; Time Taken: 1.791 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1744085536392\_0010, Tracking URL =  
[http://quickstart.cloudera:8088/proxy/application\\_1744085536392\\_0010/](http://quickstart.cloudera:8088/proxy/application_1744085536392_0010/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1744085536392\_0010

Hadoop job information for Stage-5: number of mappers: 1; number of reducers: 0

2025-05-07 07:39:48,563 Stage-5 map = 0%, reduce = 0%

2025-05-07 07:39:59,618 Stage-5 map = 100%, reduce = 0%, Cumulative CPU 3.02 sec

MapReduce Total cumulative CPU time: 3 seconds 20 msec

Ended Job = job\_1744085536392\_0010

MapReduce Jobs Launched:

Stage-Stage-5: Map: 1 Cumulative CPU: 3.02 sec HDFS Read: 8926 HDFS Write: 104 SUCCESS

Total MapReduce CPU Time Spent: 3 seconds 20 msec

OK

1	John	101	Laptop	2	60000.0	120000.0
2	Alice	102	Mouse	1	500.0	500.0
3	Bob	103	Keyboard	5	1500.0	7500.0

Time taken: 34.896 seconds, Fetched: 3 row(s)

**hive> CREATE INDEX idx\_cust\_id**

**> ON TABLE Customer\_info (Cust\_ID)**

**> AS 'COMPACT'**

**> WITH DEFERRED REBUILD;**

OK

Time taken: 0.73 seconds

**hive> SELECT**

**> SUM(o.Quantity \* i.ItemPrice) AS TotalSales,**

**> AVG(o.Quantity \* i.ItemPrice) AS AverageSales**

**> FROM**

**> Order\_info o**

**> JOIN**

**> Item\_info i ON o.ItemID = i.Item\_ID;**

Query ID = cloudera\_20250507074141\_f0342dcd-43bf-4977-a420-b0676dc697bf

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera\_20250507074141\_f0342dcd-43bf-4977-a420-b0676dc697bf.log

2025-05-07 07:41:16 Starting to launch local task to process map join;maximum memory = 932184064

2025-05-07 07:41:18 Dump the side-table for tag: 0 with group count: 3 into file:  
file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-41-11\_111\_4624799052663818332-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile20--.hashtable

2025-05-07 07:41:18 Uploaded 1 File to: file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-41-11\_111\_4624799052663818332-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile20--.hashtable (326 bytes)

2025-05-07 07:41:18 End of local task; Time Taken: 2.28 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job\_1744085536392\_0011, Tracking URL =  
[http://quickstart.cloudera:8088/proxy/application\\_1744085536392\\_0011/](http://quickstart.cloudera:8088/proxy/application_1744085536392_0011/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1744085536392\_0011

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2025-05-07 07:41:28,440 Stage-2 map = 0%, reduce = 0%

2025-05-07 07:41:37,108 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.54 sec

2025-05-07 07:41:45,544 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.06 sec

MapReduce Total cumulative CPU time: 4 seconds 60 msec

Ended Job = job\_1744085536392\_0011

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.06 sec HDFS Read: 13782 HDFS Write: 28  
SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 60 msec

OK

128000.0      42666.666666666664

Time taken: 37.219 seconds, Fetched: 1 row(s)

hive> SELECT

```
> o.OrderID,
> i.Item_Name,
> o.Quantity,
> i.ItemPrice,
> (o.Quantity * i.ItemPrice) AS TotalCost
> FROM
> Order_info o
> JOIN
> Item_info i ON o.ItemID = i.Item_ID
> ORDER BY
> TotalCost DESC
> LIMIT 1;
```

Query ID = cloudera\_20250507074242\_8f0779ab-4bae-45f2-9023-0268147739c0

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera\_20250507074242\_8f0779ab-4bae-45f2-9023-0268147739c0.log

2025-05-07 07:42:54    Starting to launch local task to process map join;maximum memory = 932184064

2025-05-07 07:42:55    Dump the side-table for tag: 0 with group count: 3 into file:  
file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-42-50\_416\_6735928633199794294-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile30--.hashtable

2025-05-07 07:42:55    Uploaded 1 File to: file:/tmp/cloudera/f6334910-0db8-4609-8414-7b60ea21ec06/hive\_2025-05-07\_07-42-50\_416\_6735928633199794294-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile30--.hashtable (329 bytes)

2025-05-07 07:42:55    End of local task; Time Taken: 1.112 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1744085536392\_0012, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1744085536392\\_0012/](http://quickstart.cloudera:8088/proxy/application_1744085536392_0012/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1744085536392\_0012

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2025-05-07 07:43:05,058 Stage-2 map = 0%, reduce = 0%

2025-05-07 07:43:12,766 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.9 sec

2025-05-07 07:43:23,219 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 4.2 sec

MapReduce Total cumulative CPU time: 4 seconds 200 msec

Ended Job = job\_1744085536392\_0012

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.2 sec HDFS Read: 13406 HDFS Write: 30  
SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 200 msec

OK

101 Laptop 2 60000.0 120000.0

Time taken: 34.954 seconds, Fetched: 1 row(s)

**hive> CREATE EXTERNAL TABLE hbase\_customer\_info (**

**> rowkey STRING,**

**> cust\_id STRING,**

**> cust\_name STRING,**

**> order\_id STRING**

**> )**

**> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'**

**> WITH SERDEPROPERTIES (**

**> "hbase.columns.mapping" = ":key,cf1:cust\_id,cf1:cust\_name,cf1:order\_id"**

**> )**

**> TBLPROPERTIES (**



```
> "hbase.table.name" = "hbase_customer_info"
```

```
> );
```

OK

Time taken: 0.828 seconds

```
hive> SELECT * FROM hbase_customer_info;
```

OK

```
row1  1      John   101
```

```
row2  2      Alice   102
```

Time taken: 0.143 seconds, Fetched: 2 row(s)

## #HBASE COMMANDS BEFORE EXTERNAL TABLE CREATION

```
[cloudera@quickstart Desktop]$ hbase shell
```

2025-05-07 07:45:12,230 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated.  
Instead, use io.native.lib.available

HBase Shell; enter 'help<RETURN>' for list of supported commands.

Type "exit<RETURN>" to leave the HBase Shell

Version 1.2.0-cdh5.13.0, rUnknown, Wed Oct 4 11:16:18 PDT 2017

```
hbase(main):001:0>
```

```
hbase(main):002:0* create 'hbase_customer_info', 'cf1'
```

0 row(s) in 2.7110 seconds

```
=> Hbase::Table - hbase_customer_info
```

```
hbase(main):003:0>
```

```
hbase(main):004:0* put 'hbase_customer_info', 'row1', 'cf1:cust_id', '1'
```

0 row(s) in 0.1810 seconds

```
hbase(main):005:0> put 'hbase_customer_info', 'row1', 'cf1:cust_name', 'John'
```

0 row(s) in 0.0130 seconds

```
hbase(main):006:0> put 'hbase_customer_info', 'row1', 'cf1:order_id', '101'
```

```
0 row(s) in 0.0090 seconds
```

```
hbase(main):007:0>
```

```
hbase(main):008:0* put 'hbase_customer_info', 'row2', 'cf1:cust_id', '2'
```

```
0 row(s) in 0.0080 seconds
```

```
hbase(main):009:0> put 'hbase_customer_info', 'row2', 'cf1:cust_name', 'Alice'
```

```
0 row(s) in 0.0060 seconds
```

```
hbase(main):010:0> put 'hbase_customer_info', 'row2', 'cf1:order_id', '102'
```

```
0 row(s) in 0.0060 seconds
```